

Class06

Anna Waters (PID: A16271985)

2024-01-25

R Functions

Functions are how we get stuff done. We call functions to do everything useful in R.

One cool thing about R is that it makes writing your own functions comparatively easy or accessible.

All functions in R have at least 3 things:

- A **Name** (we get to pick this)
- **Input Arguments** (the input to our function, one or more)
- The **body** (lines of code that do the work)

```
funname <- function(input1, input2){  
  #This body with R code  
}
```

Lets write a silly first function

```
x <- 5  
y <- 1  
x + y
```

```
[1] 6
```

```
addme <- function(x, y=1){  
  x + y  
}
```

```
addme(100,100)
```

```
[1] 200
```

```
addme(10)
```

```
[1] 11
```

Lab for today

Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped.

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

Snippet for the answers to the question on student 1

```
mean(student1) #average grade for student
```

```
[1] 98.75
```

```
mean(student2,na.rm = TRUE) # issue with NA
```

```
[1] 91
```

```
mean(student3,na.rm=TRUE) # can't just remove all NA's bc this isnt right
```

```
[1] 90
```

Come back to the NA problem.

We want to drop the lowest score before getting the `mean()`

```
#tells the lowest score but not the location
min(student1)
```

```
[1] 90
```

I found `which.min()` function. Maybe this is more useful? It should give the location of the lowest score

```
#find lowest score
which.min(student1)
```

```
[1] 8
```

Cool- it is the 8th element of the vector that has the lowest score. Can I remove this one?

```
student1[which.min(student1)]
```

```
[1] 90
```

We can use the little minus trick from indexing

```
x <- 1:5
x[-3]
```

```
[1] 1 2 4 5
```

Using the index trick to produce a vector without the lowest value and using `mean()` on the new vector to find the average without the lowest value.

```
# vector w/o the lowest score
student1[-which.min(student1)]
```

```
[1] 100 100 100 100 100 100 100
```

```
#mean without the lowest score
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Use a common shortcut and use `x` as my input

```
x <- student3
mean(x[-which.min(x)])
```

```
[1] NA
```

We still have the problem of missing values. One idea is to replace NA values with 0.

```
y <- 1:5
#replacing 3 with 1000 by seeing which = 3 & replacing
y[y == 3] <- 1000
y
```

```
[1] 1 2 1000 4 5
```

Bummer, NA's are special so == does not work

```
y <- c(1,2,NA,4,5)
y == NA
```

```
[1] NA NA NA NA NA
```

```
is.na(y)
```

```
[1] FALSE FALSE TRUE FALSE FALSE
```

How can I remove the NA elements from the vector? First I need to flip the TRUE elements to FALSE

```
#! flips logicals
!c(F,F,F)
```

```
[1] TRUE TRUE TRUE
```

```
#y[is.na(y)]
```

```
y[!is.na(y)]
```

```
[1] 1 2 4 5
```

```
y[is.na(y)] <- 1000
```

Okay lets put humpty dumpty back together again.

```
x <- student3
```

```
#assigning where NA = TRUE, the value of 0/ change NA to 0  
x[is.na(x)] <- 0  
#find and remove min value and get mean  
mean(x[-which.min(x)])
```

```
[1] 12.85714
```

Last step, now that I have a working code snippet is to make my `grade()` function

```
grade <- function(x){  
  # Change NA to 0  
  x[is.na(x)] <- 0  
  # Find and removed min value and get mean  
  mean(x[-which.min(x)])  
}
```

Testing `grade()` function on the test students

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

Read in the online gradebook

```
url <- "https://tinyurl.com/gradeinput"
gradebook <- read.csv(url, row.names = 1)

head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77

Applying the `grade()` function to the sample data from the csv.

```
# apply function uses the function on each row (margin 1) instead of needing a for loop to
results <- apply(gradebook,MARGIN = 1,FUN = grade)
results
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Q2. Using your `grade()` function and the supplied gradebook, Who is the top scoring student overall in the gradebook?

```
#which.max locates the highest scoring student
which.max(results)
```

```
student-18
18
```

```
#what the max average value is  
max(results)
```

```
[1] 94.5
```

Looking above, the top student is student 18.

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)? [2pts]

```
# apply the grade function by column to find averages per hw  
hw_results <- apply(gradebook,MARGIN = 2,mean,na.rm=T)  
  
#which homework got the lowest result  
which.min(hw_results)
```

```
hw3  
3
```

```
#lost score of the hw  
min(hw_results)
```

```
[1] 80.8
```

Mean is susceptible to outliers so `sum()` can be used instead. This combats the skew.

```
hw_sum <- apply(gradebook,MARGIN = 2,sum,na.rm=T)  
  
#which homework got the lowest result  
which.min(hw_sum)
```

```
hw2  
2
```

```
#lost score of the hw  
min(hw_sum)
```

```
[1] 1456
```

Looking above, the toughest homework, taking into account skew is hw2.

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)? [1pt]

```
# Make all NA equal to zero in mask
mask <- gradebook
mask[is.na(mask)] <- 0
#mask
```

We can use `cor()` function for correlation analysis.

```
cor(mask$hw5, results)
```

```
[1] 0.6325982
```

```
cor(mask$hw3, results)
```

```
[1] 0.3042561
```

```
apply(mask, MARGIN = 2, FUN = cor, results)
```

	hw1	hw2	hw3	hw4	hw5
	0.4250204	0.1767780	0.3042561	0.3810884	0.6325982

Looking above, hw5 is the most correlated with student results.

Q5. Make sure you save your Quarto document and can click the “Render” (or Rmarkdown”Knit”) button to generate a PDF foramt report without errors. Finally, submit your PDF to gradescope. [1pt]