

class09:Halloween Mini-Project

Anna Waters (PID: A16271985)

Here we analyze a candy dataset from the 538 website. This is a CSV file from their GitHub repository.

Importing Data

```
candy_file <- "https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-ratings.csv"

candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisp	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0

	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0		0.732		0.860	66.97	173
3 Musketeers	0	1	0		0.604		0.511	67.60	294
One dime	0	0	0		0.011		0.116	32.26	109
One quarter	0	0	0		0.011		0.511	46.11	650
Air Heads	0	0	0		0.906		0.511	52.34	146
Almond Joy	0	1	0		0.465		0.767	50.34	755

Data Exploration

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

There are 85 different types of candy in this data set.

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

There are 38 candies that are labeled as fruity in the data set.

Whats your favorite Candy

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
candy["Reese's Peanut Butter cup",]$winpercent
```

```
[1] 84.18029
```

My favorite, Reese's peanut butter cups have a win percentage of 84.18%.

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat",]$winpercent
```

```
[1] 76.7686
```

The win percent for kit kat is 76.76%

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars",]$winpercent
```

```
[1] 49.6535
```

The win percent for the tootsie roll is 49.65%.

What is the least liked candy in the dataset - lowest winpercent?

```
inds <- order(candy$winpercent)
head(candy[inds,])
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
Root Beer Barrels	0	0	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip				0	0	0	1	0.197		0.976
Boston Baked Beans				0	0	0	1	0.313		0.511
Chiclets				0	0	0	1	0.046		0.325
Super Bubble				0	0	0	0	0.162		0.116
Jawbusters				0	1	0	1	0.093		0.511
Root Beer Barrels				0	1	0	1	0.732		0.069

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744
Root Beer Barrels	29.70369

The least liked candy is the Nik L Nip with a win percent of 22.44%.

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	

numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?


The winpercent column is not in the range of 0 to 1 unlike the other columns.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

The 1 and 0 represents if the candy is chocolate(1) or not(0).

Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent)
```



Q9. Is the distribution of winpercent values symmetrical?

No, the histogram is skewed.

Q10. Is the center of the distribution above or below 50%?

The center of distribution is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
#using candy$winpercent to find all the winpercent values & taking the mean of those values
choco.mean <- mean(candy$winpercent[as.logical(candy$chocolate)])

fruity.mean <- mean(candy$winpercent[as.logical(candy$fruity)])

choco.mean
```

```
[1] 60.92153
```

```
fruity.mean
```

```
[1] 44.11974
```

On average, chocolate candy has a higher ranking than fruity candies as chocolate ranks at 60.9% and fruity at 44.1%.

Q12. Is this difference statistically significant?

```
t.test(candy$winpercent[as.logical(candy$chocolate)],
       candy$winpercent[as.logical(candy$fruity)] )
```

Welch Two Sample t-test

```
data: candy$winpercent[as.logical(candy$chocolate)] and candy$winpercent[as.logical(candy$fruity)]
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

Yes, the p.value of the t-test is smaller than 0.05 which means the null is rejected and the means are statistically significant.

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
inds <- order(candy$winpercent)
head(candy[inds,], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard bar	pluribus	sugar	percent	price	percent
Nik L Nip		0	0	0	1		0.197		0.976
Boston Baked Beans		0	0	0	1		0.313		0.511
Chiclets		0	0	0	1		0.046		0.325

Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

The five least liked are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy[order(candy$winpercent,decreasing = T),], n=5)
```


	chocolate	fruity	caramel	peanutyalmondy	nougat
Reese's Peanut Butter cup	1	0	0	1	0
Reese's Miniatures	1	0	0	1	0
Twix	1	0	1	0	0
Kit Kat	1	0	0	0	0
Snickers	1	0	1	1	1
	crispedrice	wafer	hard bar	pluribus	sugarpercent
Reese's Peanut Butter cup		0	0	0	0.720
Reese's Miniatures		0	0	0	0.034
Twix		1	0	1	0.546
Kit Kat		1	0	1	0.313
Snickers		0	0	1	0.546
	pricepercent	winpercent			
Reese's Peanut Butter cup	0.651	84.18029			
Reese's Miniatures	0.279	81.86626			
Twix	0.906	81.64291			
Kit Kat	0.511	76.76860			
Snickers	0.651	76.67378			

The top 5 most liked candies are Reese's Peanut butter cups, reese's miniatures, Twix, Kit Kats and snickers.

Q15. Make a first barplot of candy ranking based on winpercent values


```
library(ggplot2)
```

```
ggplot(candy, aes(x= winpercent, y= rownames(candy))) +  
  geom_col()
```




Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
# `reorder` is changing the rowname order in accordance with the winpercent as it is the y  
ggplot(candy, aes(winpercent, reorder(rownames(candy), winpercent))) +  
  geom_col() +  
  labs(x="Win Percent", y="")
```

```
#to fix the aspect ratio saving it can allow to manually set the ratio to the 7 by 10
ggsave('barplot1.png', width = 7, height = 10)
```




can insert any image using the exclamation mark w/ square bracket markdown syntax

Time for Color

We need a custom color vector to add color to our ggplot.

```
#setting the vector to be as long as all rows
my_cols <- rep("black", nrow(candy))
# using the as logical to assign true and false & assign chocolate color to the trues
my_cols[as.logical(candy$chocolate)] = "lightblue4"
my_cols[as.logical(candy$bar)] = "purple3"
my_cols[as.logical(candy$fruity)] = "orchid3"

ggplot(candy, aes(winpercent, reorder(rownames(candy), winpercent)))+
  geom_col(fill = my_cols) +
  labs(x="Win Percent", y="")
```



Q17. What is the worst ranked chocolate candy?


Sixlets are the worst ranked chocolate candy.

Q18. What is the best ranked fruity candy?

Starburst are the best ranked fruity candy.

[illegible]


```
ggplot(candy, aes(winpercent, pricepercent, label=rownames(candy))) +  
  geom_point(col = my_cols) +  
  geom_text()
```



```
library(ggrepel)
```

```
ggplot(candy, aes(winpercent, pricepercent, label=rownames(candy))) +  
  geom_point(col = my_cols) +  
  geom_text_repel()
```


Warning: ggrepel: 50 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Play with max.overlaps parameter in geom_text_repel()

```
ggplot(candy, aes(winpercent, pricepercent, label=rownames(candy))) +  
  geom_point(col = my_cols) +  
  geom_text_repel(max.overlaps = 5, col=my_cols)
```

Warning: ggrepel: 74 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
ord <- order(candy$winpercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Reese's Peanut Butter cup	0.651	84.18029
Reese's Miniatures	0.279	81.86626
Twix	0.906	81.64291
Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

Reese's minis have the second highest winrate of 81.86% with a pricepercent of 0.279, making it the best ranked and most economical.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord_price <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord_price,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050


The top 5 are Nik L Nip, Nestle Smarties, Ring Pops, Hershey's krackel, Hershey's milk chocolate with the least popular as Nik L Nip.

5 Exploring the correlation structure

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Fruity and chocolate values are anti-correlated with each other since the circle is large and dark red. It is unlikely that fruit and chocolate are together.

Q23. Similarly, what two variables are most positively correlated?

Other than a variable against itself, the winpercent and chocolate variables are highly correlated with a large size and dark blue color.

6 PCA


```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

```
plot(pca$x[,1:2])
```


```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])
```

```
#set the size to show the win percentage
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)
```

p




```
p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark purple), chocolate other (steel grey)",
       caption="Data from 538")
```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (dark purple), chocolate other (steel grey), f




Data from 538

##loadings plot

```
loadings <- as.data.frame(pca$rotation)

ggplot(loadings, aes(PC1, reorder(rownames(loadings), PC1))) +
  geom_col()
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard and pluribus variables are picked up in the positive PC1 direction. This makes sense because it echos what we found and what we know about candy correlations. Fruity candies are more likely to come with many like starburts or skittles and are more likely to be hard. It also makes sense they are in the opposite direction of chocolate and chocolate associated factors like nougat.